

Reflexive AI Work Group

A facilitation guide for running the first session of a facilitated work group on human-centered AI in your organization, region or online.

// WHY THIS SESSION

Build connections, clarify the initiative, and create a first broad increment of a governance framework.

Work groups meet at least three times to collaboratively build and iterate a governance framework for human-centered AI, grounded in participatory action research.

Face-to-face, ≈4 hours

Facilitated small-group work and individual reflection, ending with concrete actions to test between sessions.

5–30 participants, 11 steps

Session 1 of a three-session arc. Online is possible with breakout rooms and some modifications.

PURPOSE OF THE STUDY

Help us uncover how organizations can use AI in a human-centered, responsible manner.

A collaborative action research project run by volunteering academics and professionals, and overseen by a peer review board. Participants take part as co-researchers.

reflexiveai.org · info@reflexiveai.org

LICENSE

This work is licensed under **Creative Commons BY-NC-SA 4.0**. You may share and adapt it for non-commercial purposes, with attribution, provided any adaptations are distributed under the same license.

creativecommons.org/licenses/by-nc-sa/4.0

About this guide

Thank you for contributing to the [Reflexive AI Research Initiative](#). We are grateful for your time and effort.

This facilitation guide helps you run a **Work Group** about **reflexive AI**. You can run it in your organization, your geographical region or online. It provides a robust way to engage a broad range of participants in tackling important questions about human-centered AI, and in identifying a framework of strategies, interventions and rules to guide organizations.

The aim for each **Work Group** is to meet at least three times over a period of time to contribute to the broader [Reflexive AI Research Initiative](#). The work groups are based on **participatory action research**, in which scientists and professionals collaboratively and iteratively address societal challenges. The overall methodology is [outlined here](#).

This guide is for **Session 1** of the Work Groups. It takes **4 hours** and can be attended by 5–30 participants. We recommend a face-to-face setting, but online is possible with breakout rooms and some modifications.

Session roadmap

Each **work group** meets at least three times to collaboratively contribute to a governance framework for human-centered AI. Work groups can focus on specific themes, iterate a next version of the framework, or focus on practical translation. The broad aim for the three sessions is described below.

SESSION 1 · THIS GUIDE

- Create connections between participants.
- Clarify the aim of the initiative and how participants relate to it as co-researchers.
- Create a first, broad (and probably thin) increment of part of a framework to collect feedback on in our organizations ("What is missing?", "What is useful?").

SESSION 2

- Analyze feedback and data collected since the previous session and use it to iterate the framework.
- Reflect on the collaborative process to date and revise our Purpose-to-Practice.
- Develop interventions, strategies and actions for the themes more deeply, and test those in real organizations.

SESSION 3

- Analyze feedback and data collected since the previous session and use it to iterate the framework.
- Evaluate the process and outcomes to date (a "reality check") and identify open questions, challenges and dilemmas.
- Participants can then choose to continue with more sessions or to stop; other work groups can pick up the open questions.

The above is an outline. The aim of Sessions 2 and 3 may be adjusted based on group preferences.

Contributing to the Reflexive AI Initiative

You can run the **Work Groups** as part of our global Research Initiative, but please ensure:

- That you first contact reflexiveai.org so we can properly coordinate and ensure academic rigor, pre- and post-session surveys, and informed consent by all participants. We publish shared outcomes and materials on [GitHub](#).
- Ensure all participants are aware they are participating in a research initiative, and that the (anonymous) results from the discussions will be shared with other work groups and on GitHub.
- Avoid biasing participants by pushing them in a certain direction (i.e. "AI is great" or "AI is terrible").
- Do not modify the format unless absolutely necessary. Informed consent must always take place at the start. If you require deviations, please let us know what you changed and how.

Do you need help or have questions? Do you have feedback on this guide? Do not hesitate to contact the coordination team at info@reflexiveai.org.

How to prepare

How to invite people

- Clarify the aim of the work groups: to bring your AI challenges, experiences and dilemmas into a facilitated work group, and learn alongside peers from other organizations and academia.
- For the best outcomes, maximize the number of people with decision-making authority around AI.
- Emphasize that the discussion is part of a scientific study. Insights may be shared with [the research initiative](#).

Materials

- Print the [pre-mortem poster](#) on A0/A1 size (10 recommended).
- Print the [human-centered AI definition poster](#) (1 total).
- Print the [15% Solutions poster](#) (1 total).
- Print the [informed consent forms](#) (1 per participant).
- Print the [Study Roadmap](#) (1 total).
- A stack of empty flip charts or whiteboards.
- Bring stickies, markers and pens.
- A timer is useful to track time.
- A buzzer or tingsha bell is useful to indicate new rounds.
- Optionally, bring snacks and drinks for a relaxed atmosphere.



5 min

Step 1. Opening

We open the work group session.

1. **Welcome** the participants and thank them for their time and commitment to collaborate on this research initiative.
2. Verify that all participants have filled in the **informed consent form** that was part of the pre-workgroup survey. If somebody has not done so, they can do it now.
3. **Set expectations** that this session is about getting to know each other and the initiative, and about creating a first, broad (and probably thin) increment. We'll go deeper in the next sessions.



15 min

Step 2. Building connections

We begin with an opener to build connections, learn what brings everyone to the work group, and start thinking about what we hope to explore and achieve together.

Steps

- **5 min** Ask: *"What is your motivation for joining this work group?"* Give participants 1 minute to reflect individually. Then ask them to form pairs and discuss their responses for 4 minutes.
- **5 min** Ask participants to form new pairs. Ask: *"What do you hope we can achieve together in this work group?"* Give them 1 minute to reflect individually, followed by 4 minutes to discuss.
- **5 min** Ask participants to form new pairs again. Ask: *"What are you most curious about when it comes to human-centered AI in organizations?"* Give them 1 minute to reflect individually, followed by 4 minutes to discuss.



10 min

Step 3. Introducing the initiative

During Phase 0 of the research initiative, we created an initial Purpose-to-Practice. We will expand this with participants — acting as co-researchers — throughout the project. We present the initial Purpose-to-Practice here as a starting point.

Steps

- Briefly introduce Purpose-to-Practice and explain that we will use its five elements as a lens to introduce the initiative.
- **Purpose:** to answer the research question *“How can organizations leverage the power of artificial intelligence in a manner that is human-centered? What principles, rules, norms, interventions, processes and strategies are needed to achieve this? How can this be monitored?”*
- **Principles:**
 - **Action research** — professionals and academics iteratively develop knowledge through shared sense-making, aiming to solve real-world challenges.
 - **Co-researchers** — participants are co-researchers.
 - **Ownership** — the outcomes we create are a public good. The process is (initially) owned by the coordination team. Academic publications are owned by the writers, with all participants named as contributors.
 - **Peer review** — action research needs to be rigorous and relevant, and peer review helps protect that.
 - **Iterative** — outcomes are created incrementally, and the process is reflected on periodically with the core researchers.
- **Participants:**
 - Coordination team
 - Peer review board
 - Work groups
 - Organizations running world cafés
 - *Whoever else we feel we should include*
- **Structure:**
 - Show the initiative visualization (Study Roadmap).
- **Practices:**
 - A [public GitHub](#) as a central repository of outcomes and materials.
 - The aim of action research is to **try solutions** in real-world settings and use feedback to inform the next version — so participants can ideally test outcomes in their organizations between sessions.



🕒 20 min

Step 4. What does “human-centered AI” mean to you?

We begin by asking participants what they understand by “human-centered AI”. Rather than a theory-first definition, we use action research to create shared sense-making.

Steps

- Briefly set the stage. Explain that the goal is to create a first shared working definition of human-centered AI.
- **2 min** Ask participants to reflect silently and individually on: *“What does human-centered AI look like to you? How would you explain it to someone else?”* Invite them to **write down** a few key characteristics or principles on **index cards** or **stickies**.
- **4 min** Ask participants to pair up, exchange their ideas, and capture their understanding on **index cards** or **stickies**.
- **8 min** Ask each pair to join another pair to form groups of four. Have them compare and combine their ideas, and identify the characteristics or principles they consider most important.
- **10 min** Bring everyone back together. Collect the **stickies** or **index cards** centrally on a wall or floor for future reference. Look for common themes and differences, and use these to develop a first shared working definition of human-centered AI.
- **collect** Create pictures of the written output.



🕒 15 min

Step 5. Share research insights & initial framework

We provide a brief, neutral overview of the [scientific consensus](#) on (human-centered) AI as it relates to the research question. This grounds the discussion in scholarly conversations. We also provide a broad (empty) framework of potential themes from this literature to offer initial guidance.

Topics covered through a presentation

- A definition of AI
- AI and productivity
- AI in organizations
- Human-centered AI
- Governance of AI



10 min

Step 6. Select themes with dot voting

A governance framework covers many potential themes. We use themes from the literature as a starting point and expand and develop structure as we iterate.

In this step we invite participants to pick the theme they want to spend this first session on, based on their expertise, interest or relevance. The result will be several small groups (or individuals) working on a topic in the next part of the work group.

Potential themes for the governance framework

TO FINALIZE: The list of potential themes from literature is still being developed. The list below illustrates potential themes.

- HRM
- Leadership
- Ownership
- Work enjoyment
- Governance
- Sustainability
- Technology
- Legal

Steps

- Introduce the themes.
- **1 min** Ask participants to individually reflect on the themes and pick three they feel are most suitable for them. If a theme is missing, they can note it and suggest it in the next step.
- **2 min** Ask for themes people feel are missing. Discuss briefly and add to the list.
- **4 min** Ask participants to take a marker and distribute three dots across the themes they want to spend the rest of the session on. They can put all dots on the same theme.
- **4 min** The themes with the most dots become the focus for the remainder of this first session. Take the top 5–8 with 20+ participants, the top 3–5 with 10+ participants, and up to 3 with fewer than 10. Other themes can be tackled in future sessions or by other work groups. A participant who feels strongly about one theme can also work on it individually.
- **collect** A picture that shows the dot votes.



⌚ 40 min

Step 7. Pre-mortem

A good governance framework should prevent failure. To inform what it should address, we leverage **prospective hindsight** to identify failure states upfront. This also provides a benchmark to test the framework and interventions against (“would this scenario have been avoided if we did X?”).

We use a **pre-mortem** with a partial **TRIZ** to capture potential failure states, both likely and unlikely. These provide input for the next step. Each theme group works on its own pre-mortem poster.

Steps

- **2 min** **Round 1A** For the theme they picked, participants silently and individually **write down** responses to: *“Imagine what it would look like if AI was not human-centered for this theme. What would happen, or not happen?”*
- **8 min** **Round 1B** Invite participants at each theme to share and build on their ideas and collect them on their sheet. Encourage creativity; likely *and* unlikely outcomes can be very informative.
- **2 min** **Round 2A** Participants silently and individually **write down** responses to: *“What can we do to prevent undesirable outcomes from occurring? Or at least make them less likely?”*
- **18 min** **Round 2B** Invite participants at each theme to share and build on their ideas and collect them on their sheet. What are concrete interventions, principles, strategies and actions to prevent failure? Think broadly, but keep practical relevance.
- **2 min** **Round 3A** Participants silently and individually **write down** responses to: *“What indicators, signals or evidence can be tracked, monitored or audited to avoid undesirable outcomes?”*
- **8 min** **Round 3B** Invite participants at each theme to share and build on their ideas and collect them on their sheet. What signals, indicators and evidence can be evaluated to prevent failure, and how?



30 min

Step 8. Shift & Share

Participants focused on one theme in the previous exercise. We now invite the groups to explore what other groups created, and to use their knowledge and experience to build on those outcomes with **Shift & Share**.

Steps

- Explain the sequence of steps.
- The group at each theme picks one **presenter**. This person stays with their poster throughout the exercise.
- **5 min** The groups — except the presenter — move clockwise to the next theme. There, the presenter shares the poster with the new group and collects and records further insights.
- **20 min** Repeat at least three times, or as often as time allows.
- **5 min** Debrief with the whole group. What was the experience of creating the posters? What was surprising? What was easy? What was difficult?
- **collect** The filled-in posters.



15 min

Step 9. 15% Solutions

Action research aims to find solutions to real-world challenges. Ideally, participants take the session outcomes back to their organizations to try them out and collect feedback, which they bring into the next session. Participants can also identify other actions they want to take, or requests for help.

Steps

- **2 min** Ask participants to silently and individually reflect on: *"Based on what we created here today, what is a step you can take to test it out in your organization or collect feedback on it? What is within your mandate and available resources? Where do you need help from others in this room?"*
- **8 min** Invite people to form small groups to share their actions and coach each other to make them more concrete and doable. Capture actions on **index cards** or **stickies**.
- **5 min** Ask participants who feel comfortable to share their actions, along with any help requests they have.
- **collect** The stickies or index cards with actions.



5 min

Step 10. What's next

Close the session by thanking everyone for their time. If you opted to participate in the Research Initiative, reiterate that the anonymous results will be shared with the initiative, and that research updates will be available on reflexiveai.org.

Provide a brief preview of the coming sessions:

- In Session 2, we bring back feedback from our organizations and use that — and other data — to go deeper.
- In Session 3, we iterate again and reflect on the outcomes and the process to achieve them.

We recommend going through this checklist:

- **collect** Filled-in [informed consent forms](#) for every participant who did not submit one upfront (not needed if you won't contribute to the research initiative).
- **collect** The filled-in poster for the definition of human-centered AI.
- **collect** The filled-in poster for the 15% Solutions.
- **collect** The filled-in posters for the various themes.
- **collect** A picture that shows the dot votes for the themes.



Step 11. Share the results

When you've completed the session, please share the following with the Research Initiative if you opted to do so.

- Clear pictures or scans of the posters, including any stickies or content on them. Please remove or blur people's names and company-sensitive information. These will be published on [our public GitHub](#) to benefit the initiative and other work groups.
- A brief review by you and any other facilitators of how the session went. What worked? What did not? Where did you observe tension or friction?
- Scans or pictures of filled-in [informed consent forms](#) for all participants. Unfortunately, we cannot process *any* data without informed consent.
- Send the materials to info@reflexiveai.org.



Post-session activities

After the session, facilitators individually capture their experiences, observations and reflections in a [personal log](#). The aim is to document what was noticed — surprises, tensions, assumptions and emerging questions — and how this informs the next sessions.

If the sessions contribute to the research initiative, a short follow-up survey will be sent to the participants' email addresses.