



Reflexive AI Work Group #1

A facilitation guide for running the first session of a facilitated work group on human-centered AI in your organization, region or online.

// AIM FOR SESSION #1

Build connections, clarify the initiative, and create a first broad increment of a governance framework.

Face-to-face, ≈3 hours

Facilitated small-group work and individual reflection, ending with concrete actions to test between sessions.

5–30 participants, 11 steps

Session 1 of a three-session arc. Online is possible with breakout rooms and some modifications.

PURPOSE OF THE STUDY

Help us uncover how organizations can use AI in a human-centered manner.

A collaborative action research project run by volunteering academics and professionals, and overseen by a peer review board. Participants take part as co-researchers.

reflexiveai.org · info@reflexiveai.org

LICENSE

This work is licensed under **Creative Commons BY-NC-SA 4.0**. You may share and adapt it for non-commercial purposes, with attribution, provided any adaptations are distributed under the same license.

creativecommons.org/licenses/by-nc-sa/4.0

About this guide

Thank you for contributing to the [Reflexive AI research initiative](#). We are grateful for your time and effort.

This facilitation guide helps you run a **work group** about **reflexive AI**. Reflexivity means *"to step back and view oneself from a distance to become critically aware of oneself in relation to larger contexts"*. For this initiative, it means reflecting on how your use of AI is (or is not) influenced by your values, principles and norms and how this in turn influences the people around you, your organization and broader society.

You can run a work group in your organization, your geographical region or online. It provides a robust way to engage a broad range of participants in tackling important questions about human-centered AI, and in identifying a framework of strategies, interventions and rules to guide organizations.

The aim for each **work group** is to meet at least three times over a period of time to contribute to the broader [Reflexive AI research initiative](#). The work groups are based on **participatory action research**, in which scientists and professionals collaboratively and iteratively address societal challenges. Work groups iterate the next version of the framework, or focus on practical translation. They can focus on specific themes. The broad aim for the three sessions is described below. The methodology is [outlined here](#).

This guide is for **Session 1**. It takes **3 hours** and can be attended by 5–30 participants. We recommend a face-to-face setting, but online is possible with breakout rooms and some modifications.

Session roadmap

SESSION 1 · THIS GUIDE

- Create connections between participants.
- Clarify the aim of the initiative and how participants relate to it as co-researchers.
- Create a first, broad increment of (part of) the framework.
- Participants plan a concrete intervention they want to try within their organization before the next session, for a segment of the framework and at a scale manageable for them.

SESSION 2

- Evaluate actions taken since the previous session, as well as input from the World Cafés, and use it to iterate the framework.
- Reflect on the collaborative process to date and revise our overall approach (Purpose-to-Practice).
- Develop interventions, strategies and actions for the themes more deeply, and test those in real organizations.
- Participants plan a concrete intervention they want to try within their organization before the next session, for a segment of the framework.

SESSION 3

- Evaluate actions taken since the previous session, as well as input from the World Cafés, and use it to iterate the framework.
- Evaluate the process and outcomes to date (a "reality check") and identify open questions, challenges and dilemmas.
- Participants can then choose to continue with more sessions or to stop; other work groups can pick up the open questions.

The above is an outline. The aim of Sessions 2 and 3 may be adjusted based on group preferences.

Contributing to the Reflexive AI initiative

We designed this guide for work groups that contribute to our global research initiative. This is not required — you can run this work group within your organization in isolation. If you do contribute to the research initiative, please ensure:

- That you **coordinate with our team** at reflexiveai.org so we can ensure academic rigor, pre- and post-session surveys, and informed consent by all participants. We publish outcomes and materials on [GitHub](#).
- That **all participants are aware they are participating in a research initiative**, and that the (anonymous) results from the discussions will be shared with other work groups and on GitHub.
- That you do **not bias participants** by pushing them in a certain direction (e.g. “AI is great” or “AI is terrible”).
- That you do **not modify the format** unless absolutely necessary. Informed consent must always take place at the start. If you require deviations, please let us know what you changed and how.

Please contact us at info@reflexiveai.org with questions.

How to prepare

How to invite people

- Clarify the aim of the work groups: to bring your AI challenges, experiences and dilemmas into a facilitated work group, and to learn alongside peers from other organizations and academia.
- For the best outcomes, maximize the number of people with decision-making authority around AI.
- Emphasize that the discussion is part of a scientific study. Insights may be shared with [the research initiative](#).

Materials

- Print the [pre-mortem poster](#) on A0/A1 size (10 recommended).
- Print the [human-centered AI definition poster](#) (1 total).
- Print the [15% Solutions poster](#) (1 total).
- Print the [informed consent forms](#) (1 per participant).
- Print the [working definition for human-centered AI](#) (1 total).
- Print the [Study Roadmap](#) (1 total).
- Print the [Purpose-to-Practice Worksheet](#) (1 total). Populate it with pre-written stickies beforehand to describe the elements from Step 3 in this guide, so as to illustrate that this is a living document, not a static given.
- Optional: print the [Helpful Resources poster](#) on A0/A1 size (several to spread around the room).
- Optional: print the [Give & Get Help poster](#) on A0/A1 size (several to spread around the room).
- A stack of empty flip charts or whiteboards.
- Bring stickies, markers and pens.
- A timer is useful to track time.
- A buzzer or tingsha bell is useful to indicate new rounds.
- Optionally, bring snacks and drinks for a relaxed atmosphere.



5 min

Step 1. Opening

We open the first work group session.

1. **Welcome** the participants and thank them for their time and commitment to collaborate on this research initiative.
2. Verify that all participants have filled in the **informed consent form**. If somebody has not done so, they can do it now.
3. **Set expectations** that this session is about getting to know each other and the initiative, and about creating a first, broad (and probably thin) increment. We'll go deeper in the next sessions.



15 min

Step 2. Build connections with Impromptu Networking

We begin with an opener to build connections, learn what brings everyone to the work group, and start thinking about what we hope to explore and achieve together.

Steps

- **5 min** Ask: *"What is your motivation for joining this work group?"* Give participants 1 minute to reflect individually. Then ask them to form pairs and discuss their responses for 4 minutes.
- **5 min** Ask participants to form new pairs. Ask: *"What do you hope we can achieve together in this work group?"* Give them 1 minute to reflect individually, followed by 4 minutes to discuss.
- **5 min** Ask participants to form new pairs again. Ask: *"What are you most curious about when it comes to human-centered AI in organizations?"* Give them 1 minute to reflect individually, followed by 4 minutes to discuss.



10 min

Step 3. Introducing the initiative and its underlying approach with Purpose-to-Practice

During the preparations of the research initiative, we created an initial [Purpose-to-Practice](#). A Purpose-to-Practice (P2P) is a meta practice where the five essential elements for an initiative are defined collaboratively: purpose, principles, participants, structure and practices. Over time, this will be expanded with participants — acting as co-researchers — throughout the project. We present the initial Purpose-to-Practice as a starting point. The elements below should already be visible on the P2P with pre-written stickies. This emphasizes that the P2P changes as we learn more.

Steps

- Briefly introduce Purpose-to-Practice and explain that we will use its five elements as a lens to introduce the initiative as an enduring initiative.
- **Purpose:**
 - To answer the **research question** *“How can organizations leverage the power of artificial intelligence in a manner that is human-centered? What principles, rules, norms, interventions, processes and strategies are needed to achieve this? How can this be monitored?”*
 - To contribute to **scientific knowledge and societal awareness** around human-centered AI.
 - To **try potential solutions** in real-world settings to learn what else is needed.
- **Principles:**
 - **Action research** — professionals and academics iteratively develop knowledge and take action through shared sense-making, aiming to solve real-world challenges.
 - **Co-researchers** — participants are co-researchers.
 - **Ownership** — the outcomes we create are a public good. The process is (initially) owned by the coordination team, but we happily invite participants to co-own this. Academic publications are owned by the writers, with all participants named as contributors.
 - **Peer review** — action research needs to be rigorous and relevant, and peer review by our board helps protect that.
 - **Iterative** — outcomes are created incrementally, and the process is reflected on periodically with the co-researchers.
- **Participants:**
 - Coordination team
 - Peer review board
 - Work groups (co-researchers)
 - Organizations running World Cafés
 - *Whoever else we feel we should include*
- **Structure:**
 - Show the initiative visualization ([Study Roadmap](#)).
- **Practices:**
 - A [public GitHub](#) as a central repository for materials, including the framework itself and guidance to apply it.
 - Periodic collaborative reflection on the process of this initiative (Purpose-to-Practice) to identify improvements.
 - Academic publications in peer-reviewed journals.



🕒 20 min

Step 4. Discover what “human-centered AI” means with 1-2-4-ALL

We begin by asking participants how they understand “human-centered AI”. Rather than a theory-first definition, we use action research to build a working definition through shared sense-making.

Steps

- Set the stage. Explain that the goal is to create a first working definition of human-centered AI.
- **2 min** Ask participants to reflect silently and individually on: “*What does human-centered AI look like to you? How would you explain it to someone else?*” Invite them to **write down** a few key characteristics, principles or a definition on **index cards** or **stickies**.
- **4 min** Ask participants to pair up, exchange their ideas, and capture their understanding on **index cards** or **stickies**.
- **8 min** Ask each pair to join another pair to form groups of four. Have them compare and combine their ideas, and identify the characteristics or principles they consider most important.
- **5 min** Bring everyone back together. Collect the **stickies** or **index cards** centrally on a wall or floor for future reference. Look for common themes and differences, and use these to develop a first shared working definition of human-centered AI.
- **1 min** Clarify that one aim of this initiative is to develop a good definition for human-centered AI. Present the **incomplete working definition (PDF)** for human-centered AI: “*creating and using AI to help people perform their work in a more efficient way (faster and more accurately), while making sure that they remain in control, feel motivated, stay connected socially to their colleagues, and understand how AI affects them with autonomy to question, change, or reject its decisions*”.
- **collect** Create pictures of the written output.



🕒 15 min

Step 5. Brief presentation of research on human-centered AI & initial framework

We provide a brief, neutral overview of the [scientific consensus](#) on (human-centered) AI as it relates to the research question. This grounds the discussion in scholarly conversations. We also provide a broad (empty) framework of potential themes from the literature to offer initial guidance.

Topics covered through a presentation

- A definition of AI
- AI and productivity
- AI in organizations
- Human-centered AI
- Governance of AI



10 min

Step 6. Determine focus for this work group through dot voting

A governance framework covers many potential themes. We use themes from the literature as a starting point and expand and develop the structure as we iterate.

In this step, the group decides which themes it will focus on in this work group. This allows groups to focus on those areas where their expertise, interest and experience lies. Other themes may be picked up by other work groups.

Initial themes for the governance framework

We provide an initial set of themes based on [Tavantzis et al. \(2027\)](#) and interviews with experts performed by Z. Kaskin (2026), to be expanded and refined by participants.

- Purpose & Strategy
- Human Agency
- Governance & Ethics
- Review & Adaptation
- Meaningful Work
- Work Design
- Workforce & Learning
- Quality & Oversight
- Communication
- Participation
- Questioning Culture
- Power & Dynamics
- Risk & Boundaries
- Legal & Compliance
- IP & Data Ownership
- Sustainability

Steps

- Introduce the themes.
- **1 min** Ask participants to individually reflect on the themes and pick three they feel are most suitable for them to work on. If a theme is missing, they can note it and suggest it in the next step.
- **2 min** Ask for themes people feel are missing. Discuss briefly and add to the list.
- **4 min** Ask participants to take a marker and distribute three dots across the themes they want to spend their effort on in this and the coming sessions. They can put multiple dots on the same theme.
- **4 min** The themes with the most dots become the focus for this work group. Take the top 5–8 with 20+ participants, the top 3–5 with 10+ participants, and up to 3 with fewer than 10. Other themes can be addressed by other work groups. A participant who feels strongly about one theme can also work on it individually.
- **collect** A picture that shows the dot votes.



40 min

Step 7. Create an initial framework with Pre-Mortem and Shift & Share

A good governance framework should prevent failure. To inform what a framework for human-centered AI should address, we leverage **prospective hindsight** to identify failure states upfront. Prospective hindsight is a cognitive strategy that involves imagining potential future failures or challenges before they occur, allowing individuals to identify risks and take preventive measures. This will also provide a benchmark against which the framework and interventions can be tested later ("would this scenario have been avoided if we did X?").

We start with a **pre-mortem** to capture potential failure states, both likely and unlikely. This provides an angle to define strategies/interventions/actions to prevent such failures, and the measures to monitor this. Each theme group works on its own pre-mortem poster.

Steps

- **2 min** **Round 1A** For the theme they picked, participants silently and individually **write down** responses to: *"Imagine what it would look like if AI was not human-centered for this theme. What would happen, or not happen?"*
- **8 min** **Round 1B** Invite participants at each theme to share and build on their ideas and collect them on their sheet. Encourage creativity; likely *and* unlikely outcomes can be very informative.
- **2 min** **Round 2A** Participants silently and individually **write down** responses to: *"What can we do to prevent undesirable outcomes from occurring and encourage desirable outcomes to happen?"*
- **18 min** **Round 2B** Invite participants at each theme to share and build on their ideas and collect them on their sheet. What are concrete interventions, principles, strategies and actions to prevent failure? Who should perform them? Think broadly, but keep practical relevance.
- **2 min** **Round 3A** Participants silently and individually **write down** responses to: *"What indicators, signals, or evidence can be tracked, monitored, or audited to avoid undesirable outcomes and encourage desirable ones?"*
- **8 min** **Round 3B** Invite participants at each theme to share and build on their ideas and collect them on their sheet. What signals, indicators and evidence can be evaluated to prevent failure, and how? Who performs these evaluations?



30 min

Step 8. Shift & Share

Participants focused on one theme in the previous exercise. We now invite the groups to explore what other groups created, and to use their knowledge and experience to build on those outcomes with **Shift & Share**.

Steps

- Explain the sequence of steps.
- The group at each theme picks one **presenter**. This person stays with their poster throughout the exercise. If you run this work group with fewer than 8 participants, you can also tour the posters as a whole group.
- **5 min** The groups — except the presenter — move clockwise to the next theme. There, the presenter shares the poster with the new group and collects and records further insights.
- **20 min** Repeat at least three times, or as often as time allows.
- **5 min** Debrief with the whole group. What was the experience of creating the posters? What was surprising? What was easy? What was difficult?
- **collect** The filled-in posters.



15 min

Step 9. Identify actions with 15% Solutions

Action research aims to implement solutions to real-world challenges. We use the Liberating Structure **15% Solutions** to help participants think about this. Ideally, participants develop an intervention to try within their own organization based on some part of the framework at a scale that is manageable (i.e. work group, team, department, organization). By taking action, participants generate useful feedback to inform future actions and further develop the framework. Examples include organizing an internal workshop to raise awareness, running a Reflexive AI World Café, changing a procedure, introducing policy, involving colleagues to develop an intervention, launching an approach, starting a Community of Practice, changing a workflow, collecting data, implementing strategy or some other concrete action.

Steps

- **4 min** Ask participants to silently and individually reflect on: *“Based on what we created today, what is an intervention you can implement within your organization to test, develop or deepen the framework before the next session? What is within your mandate and available resources?”*
- **6 min** Invite people to form pairs to share their actions and coach each other to make them more concrete and doable. Capture actions on **index cards** or **stickies**.
- **5 min** Ask participants who feel comfortable to share their actions.
- **collect** The stickies or index cards with actions.



5 min

Step 10. What's next?

Close the session by thanking everyone for their time. If this group participates in the research initiative, reiterate that the anonymous results will be shared with the initiative, and that research updates will be available on reflexiveai.org.

Provide a brief preview of the coming sessions:

- In Session 2, collect and analyze feedback and other data to iterate the framework and go deeper.
- In Session 3, we iterate again and reflect on the outcomes and the process to achieve them.

We recommend going through this checklist:

- **collect** Filled-in [informed consent forms](#) for every participant who did not submit one upfront (not needed if you won't contribute to the research initiative).
- **collect** The filled-in poster for the definition of human-centered AI, including any stickies. This poster should be put up in the next sessions to let learnings physically travel with the group as "external memory".
- **collect** The filled-in poster for the 15% Solutions, including any stickies. This poster should be put up in the next sessions to let learnings physically travel with the group as "external memory".
- **collect** The filled-in posters for the various themes, including any stickies. These posters should be put up in the next sessions to let learnings physically travel with the group as "external memory".
- **collect** A picture that shows the dot votes for the themes.



Step 11. Share the results

When you've completed the session, please share the following with the research initiative if you opted to do so.

- Clear pictures or scans of the posters, including any stickies or content on them. Please remove or blur people's names and company-sensitive information. These will be published on [our public GitHub](#) to benefit the initiative and other work groups.
- A brief review by you and any other facilitators of how the session went. What worked? What did not? Where did you observe tension or friction?
- Scans or pictures of filled-in [informed consent forms](#) for all participants. Unfortunately, we cannot process *any* data without informed consent.
- Send the materials to info@reflexiveai.org.



Post-session activities

After the session, facilitators individually capture their experiences, observations and reflections in a [personal log](#). The aim is to document what was noticed — surprises, tensions, assumptions and emerging questions — and how this informs the next sessions.

If the sessions contribute to the research initiative, a short follow-up survey will be sent to the participants' email addresses.